

# Training and Testing an Autonomous Driving System

Michael Martinez

## ADS Background and Project Overview

Autonomous Driving Systems (ADS) have often been a phenomenon confined to science fiction media. However, recent advancements in computer vision, computation power, and big data technologies have allowed this phenomenon to become a reality. At the heart of every ADS or self-driving car is the ability to sense its environment. In fact, this ability is the most important feature in driving an automobile. We humans almost take this ability for granted when we get behind the wheel, but for an ADS the process of learning this task and implementing it takes years of optimization and incorporation of several technologies.

This project aims to optimize a very important variable that an ADS must evaluate after analyzing its environment: the steering angle. The steering angle is incredibly important to a self-driving car to ensure that it handles turns well, stays in its lane, avoids objects and collisions, and performs correctly in several other scenarios.

Udacity has open sourced software that puts a user behind the wheel of a car. The user has the choice of two tracks to drive around: a lake or a mountain. This simulator, built on Unity, has a couple of features key to this project - the ability to *collect* various driving data, and the ability to *test* a trained model on a specified track. The simulated car is equipped with 3 cameras on its windshield that capture a center, left, and right view of the car's frontal surroundings. Furthermore, the car is equipped with sensors that measure throttle, speed, reverse, and steering angles. Thus, the data collected consists of 3 images and 4 measured variables per recorded observation. To generate this data, a user drives around the track with the *record* option toggled on. Therefore, my project will be an exercise in *behavioral cloning*- where I generate my own data and train a model based upon it.

Behavioral cloning is a specification of artificial intelligence. In this regime, an AI learns from a *human's* interpretation of a certain task. In other words, an AI seeks to *clone* a human's *behavior* as they accomplish a specific task. In my instance, the model I train will attempt to clone the behavior I exhibit as I drive around the track. This type of AI specification is very popular in today's leading real self-driving cars. And this makes sense because after all, driving was first made to accommodate humans and therefore the infrastructure, laws, and decisions have all been installed with humans in mind. It is advantageous then, for an ADS to clone a human's driving behavior.

Once sufficient data is collected, it will be imported into a Jupyter notebook to be analyzed and processed. For this project, I will be using a leading convolutional neural network (CNN) design developed by NVIDIA and implementing it using Keras. As with any CNN, images usually must be preprocessed to ensure optimal model performance. Therefore, I will preprocess all images according to industry-standard ADS techniques before training a model on them.

After training a model, the final step is to put it to the test on the track. The car will be trained on the lake track *only*, which means it will have never have “seen” the mountain track. Due to this training constraint, I fully expect the trained model to be able to complete several laps of the lake track without going off track or colliding into anything. However, the lake and mountain tracks differ quite dramatically in several aspects such as surrounding environment, turn angles, inclines, brightness, and several more variables. Therefore, the mountain track will serve as a nice benchmark to see how well the model can generalize to new environments.

This project aims to serve as a foundational platform for an autonomous driving system by addressing a fundamental variable- steering angle. Ultimately the results of this project seek to benefit the self-driving car industry and contribute to ADS technologies.

## Gathering, Analyzing and Processing Data

As described in the previous section, I was responsible to generate and gather my own data in the simulation. I gathered data exclusively on the lake track to train my model and saved the additional track as a validation track. The two tracks as seen in the simulation are shown below in Fig. 1.



*Figure 1. Screen Captures of the simulated car on the lake (shown on the left) and mountain (shown on the right) tracks. The lake track serves as the training track, and the mountain track serves as the validation track.*

There were several iterations of the generating process that were performed based upon the data analysis and model performance. The final iteration of the gathering process included 15759 observations, each of which including the center, left, and right frontal images and the 4 additional sensory variables. This final iteration included several laps around the lake track in *both* directions. The lake track is heavily left-turn biased and therefore not an ideal model to train on. Therefore, I included data driven in the reverse direction of the track, thus countering the

left-turn bias with right-turn data. Furthermore, I padded the data with supplemental *recovery* data to optimize the cars ability to drive in the center of the lane. I define a recovery as the process of the car steering back into the center of the lane after it veers onto the edge of the lane. I load and format the gathered data into a pandas dataframe to begin the analysis. The dataframe consists of file paths to the images (center, left, and right) on my local disk as well as the 4 additional metrics- speed, throttle, reverse, and steering angle. The camera views are shown below in Fig. 2.

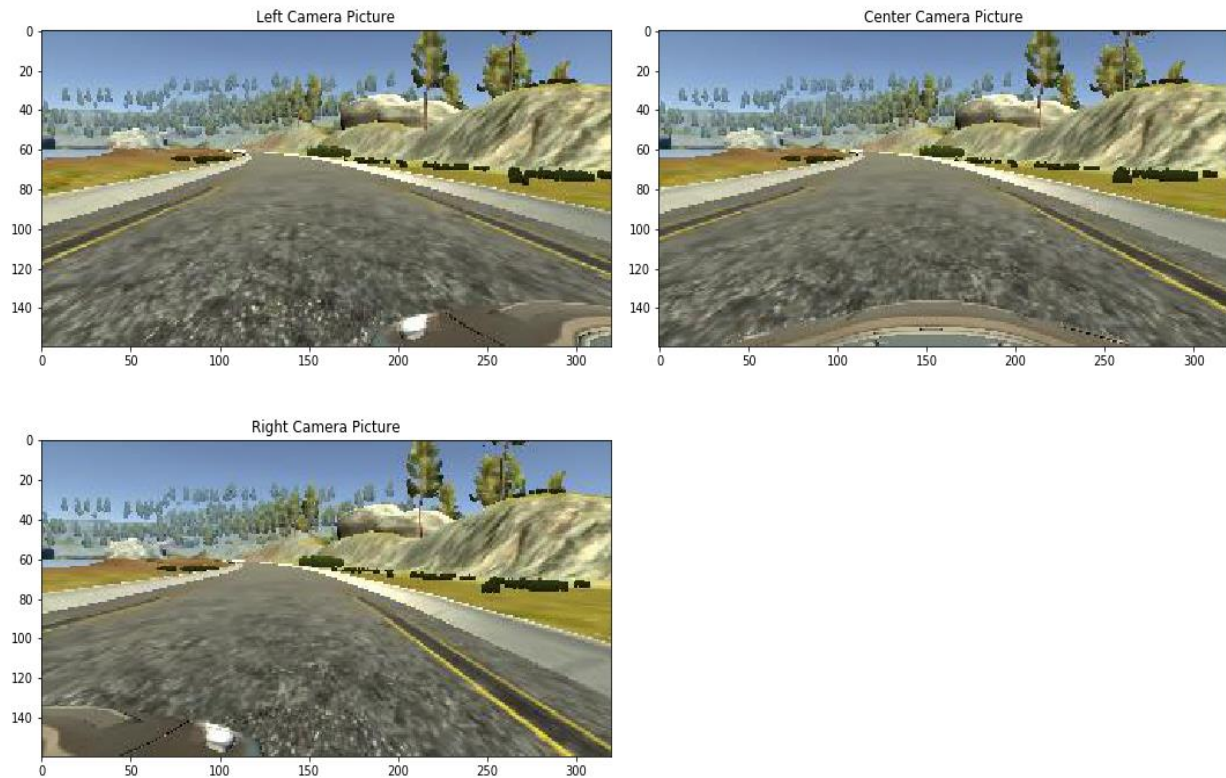


Figure 2. Image captures from the car's cameras. The car has a left, center, and right camera installed on its windshield.

Again, I am only concerned with the steering angle so I perform an initial visual analysis to understand the data distribution shown in Fig. 3.




Figure 3. Distribution of the steering angles as recorded by the training laps.

From this steering angle histogram, we can clearly see an overwhelming small angle bias. This immediately raises a red flag. While the lake track can be inferred to be heavily straight from the data and therefore only need small steering angles, there are a few turns which necessitate steering angles of magnitude .25 and greater. This can also be seen from the distribution. If I let the model train on this distribution alone, it will most certainly overfit on small angles and not be able to handle the sharper turns. I therefore define a threshold defined by the orange bar. I define a function to iterate through the bins of the data and keep a *maximum* of 750 data points *per* bin. Importantly, I shuffle the data in each bin such that the whole track still remains represented and described by the steering angle data. This result of this function then cuts the small angle bin to a more reasonable proportion of data, while still representing the whole track. The effect of this function is shown below in Fig. 4.




Figure 4. The distribution of steering angles after applying a cutoff function to achieve a more balanced distribution.

After cutting the data by the defined threshold, the resulting size is 5745 observations. The final step in this process then, is to split the data into training and validation data. This is done using the Sci-kit learn library. I initialize the train-test split as 80-20 (Fig. 5), meaning I allow my model to train on 80% of the data, and then validate on the remaining 20%. In plain numbers, this means 4596 observations to train on, and 1149 observations to test on.

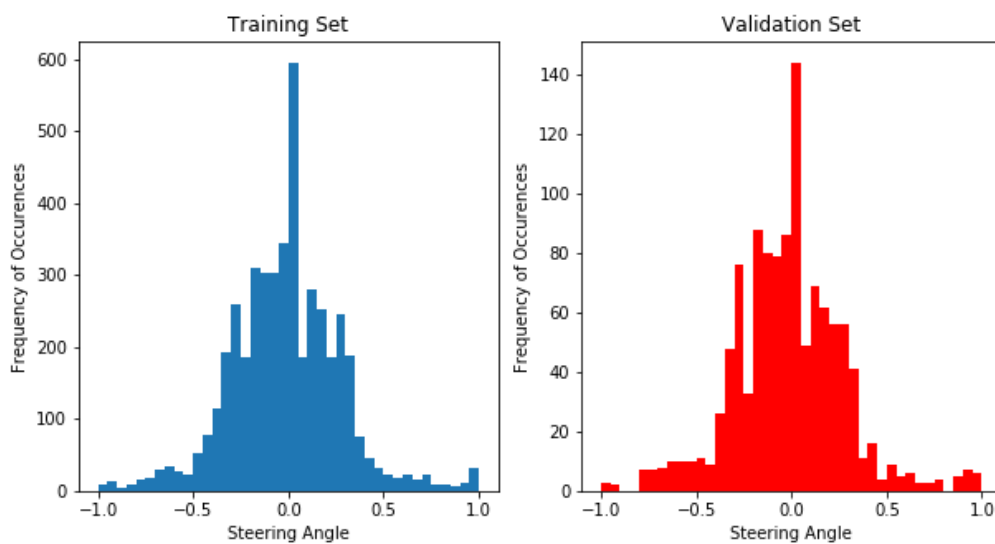


Figure 5. Distribution of steering angle data split into the training and validation sets

